TTP – Time-Triggered Protocol for Fault-Tolerant Real-Time Systems

Milestone 3: Critical Evaluation and Transfer

Jaleel Ur Rehman
Electronics Engineering
Hochschule Hamm-Lippstadt
jaleel.ur-rehman@stud.hshl.de

Concrete Strengths

- **Temporal encapsulation**:
  - Each node's behavior is fully specified by its TDMA slot boundaries.
  - A node can be developed, tested, and verified independently from all other nodes in both the value and time domains.
  - Local timing changes do not propagate system-wide – no system-wide regression testing is required after modifying one node.
  - In contrast, ET architectures have global timing side effects: a local task modification can shift message priorities and invalidate the entire system's worst-case timing analysis.
- **Replica determinism**:
  - All FTU replicas receive the same inputs at the same global time and perform the same sequence of state transitions.
  - Outputs of replicas are identical without any additional voting protocol – transparent active redundancy is architecturally straightforward.
  - ET architectures do not provide a solution to replica determinism; replicas may diverge due to differing event arrival orders.
- **Integrated services without overhead**:
  - Clock synchronization, membership tracking, group acknowledgment, and redundancy management are provided within the existing TDMA frame structure.
  - No extra protocol messages are added; these services come at zero bandwidth cost.
- **High data efficiency (53.3%)**:

- More than double the efficiency of CAN (25.3%), J1850 (27.1%).
- Structural advantage from a-priori message knowledge - not a tunable parameter.
- Directly translates into lower hardware bandwidth requirements for the same message set.

- **Load-independent worst-case timing**:
  - TDMA guarantees a fixed upper bound on message latency regardless of instantaneous bus load.
  - This is a hard requirement for closed-loop safety-critical control (e.g., brake-by-wire actuator update at 1 ms period).
- **Trivial schedulability analysis**:
  - Static schedule means schedulability is provable by inspection of the TDMA table.
  - No rate-monotonic, deadline-monotonic, or response-time analysis needed.
  - Simplifies safety case construction for IEC 61508 / ISO 26262 / DO-178C certification.

Concrete Limitations and Risks

- **Fixed schedule inflexibility**:
  - Adding or removing any FTU node requires creating a new mode definition and redeploying the updated schedule to *all* nodes in the ensemble.
  - In a production vehicle with 50+ ECUs, this requires a full software update campaign. No hot-plug capability.
  - Although mode changes handle predefined reconfigurations, unplanned topology changes are not supported.
- **Bandwidth waste under low or irregular load**:
  - Statically allocated TDMA slots remain idle when a node has no data to send.
  - Systems with bursty traffic (sensor fusion, event logs, diagnostic dumps) permanently reserve bandwidth they rarely use.
  - On a 250 kbit/s bus, over-provisioning is especially costly.
- **Fail-silent assumption is hardware-expensive and potentially incomplete**:
  - Enforcing fail-silence requires dual-redundant hardware, dedicated watchdog circuits, and comprehensive self-test coverage that detects all internal faults before any output is produced.
  - Implementation cost is significant; prohibitive for low-cost nodes.
  - A node with a fault in its own self-test logic may produce incorrect output while *believing* it has shut down correctly – the fail-silent guarantee then silently breaks.

- **Byzantine faults are completely excluded**:
  - If a node transmits inconsistent values to two receivers in the same TDMA slot (a physically possible fault mode for a node with a partially failed output driver), TTP's membership and C-state mechanisms do not detect the inconsistency.
  - This is a known, documented gap in TTP's failure model.
- **Clock synchronization scalability**:
  - The fault-tolerant averaging algorithm requires that a node can observe a majority of transmissions before computing a correction.
  - Over long physical distances, propagation delay variation causes clock offset to grow. The granule g must be enlarged to compensate, reducing effective bandwidth.
  - Large, geographically distributed ensembles (e.g., a factory floor) may exceed practical synchronization accuracy.
- **Cold-start collision risk**:
  - Simultaneous power-up (e.g., global reset after power failure) has all nodes contending to transmit I-frames.
  - The probabilistic backoff (I_1, I_3 waiting intervals) reduces collision probability but provides no formal upper bound on start-up time under worst-case simultaneous restart.

Assumptions That May Be Unrealistic in Modern Systems

- **Periodic workloads only**:
  - TTP is designed for periodic communication. Sporadic or aperiodic messaging is handled only via predefined mode changes.
  - Autonomous driving generates sensor data at rates that vary with environmental complexity – dense urban scenarios produce far more LiDAR/camera data than a highway. This variability is not statically schedulable.
- **Static, fully known topology at design time**:
  - Modern automotive software-defined vehicles support OTA updates that add new software functions and may require new ECU communication patterns at runtime.
  - Industrial IoT systems add, replace, and reconfigure sensors continuously.
  - TTP's design-time topology constraint is incompatible with open or dynamically reconfigurable systems.
- **No failure occurs during a mode change**:
  - The membership and mode-change protocols depend on this assumption.

- Under a fully asynchronous fault model this cannot be guaranteed: a node failure arriving mid-mode-change could cause a C-state divergence that the protocol does not handle cleanly.

- **Homogeneous hardware cost across all nodes**:
  - The fail-silent assumption implicitly requires all nodes to afford redundant hardware. In body electronics (window control, seat adjustment, interior lighting), this cost is not economically justifiable.

Applicability to Realistic Embedded Systems Scenarios

Scenarios Where TTP Is Suitable

- **Brake-by-wire / steer-by-wire**: Periodic actuation commands at fixed rates (typically 1–10 ms), hard deadlines, ASIL-D safety integrity, static topology determined by vehicle platform. Exactly TTP's design target.

- **Aircraft flight control**: Fixed-rate inertial and air data sensor sampling, static airframe topology per aircraft type, long certification lifecycle that benefits from formal timing proofs.

- **Railway signaling and interlocking**: Periodic safety-state updates, static track infrastructure, high safety integrity requirements with decades-long operational lifecycles.

- **Industrial robot joint control**: Fixed-rate torque/position loops (1 kHz), static controller-to-actuator wiring, determinism required for coordinated multi-axis motion without jitter.

- **Engine management (periodic sensor sampling)**: Temperature, pressure, and crankshaft position are sampled at fixed intervals; TTP's TDMA model aligns naturally with this.

Scenarios Where TTP Is Unsuitable

- **Autonomous driving perception**: Camera and LiDAR bandwidth is variable and far exceeds 250 kbit/s; OTA updates change communication patterns dynamically.

- **In-vehicle infotainment**: Bursty video/audio traffic, no hard real-time requirements, cost-sensitive. Automotive Ethernet (TSN) or CAN-FD is more appropriate.

- **Industrial IoT with dynamic topology**: Frequent sensor addition/replacement makes schedule redeployment operationally impractical.

- **Low-cost body electronics**: Window controls, seat memory, interior lighting. CAN's simplicity and $< \$1$ silicon cost per node are decisive; fail-silent hardware redundancy cannot be justified.

- **Over-the-air firmware update channels**: Unpredictable large-data transfers conflict with static TDMA allocation.

Runtime, Memory, and Scalability Implications

- Bus bandwidth: 250 kbit/s is sufficient for periodic control loops with small message payloads. Camera/LiDAR/radar streams require 100×–10000× more bandwidth. TTP cannot serve as a sensor bus in modern ADAS systems.
- Schedule storage: grows linearly with ensemble size. For a large ensemble (50+ FTUs), schedule storage in non-volatile memory may be a constraint on μC-class devices.
- TDMA round length grows with ensemble size, increasing the worst-case message latency for all nodes. There is a fundamental tension between ensemble scalability and tight control loop deadlines.
- Schedule deployment: n nodes each require a full schedule update per topology change. Deployment time grows linearly; update windows must be scheduled into maintenance downtime.

Possible Extensions and Open Research Questions

- **Dynamic schedule reconfiguration**: Can a TTP-like protocol support lightweight partial schedule updates (e.g., adding one FTU slot) without full redeployment, while preserving temporal encapsulation? Hot-patching TDMA schedules is an open design challenge.
- **Mixed-criticality integration**: TTP assigns equal static priority to all scheduled slots. Mixed-criticality scheduling models allow high-criticality tasks to borrow bandwidth from low-criticality tasks under overload – integrating this into a TDMA bus protocol is an open problem.
- **Byzantine fault tolerance at low cost**: Is there a hybrid architecture that detects inconsistent node transmissions using TTP's existing dual-channel redundancy, without requiring the full $3f + 1$ node count of classical BFT?
- **TTP principles on high-speed Ethernet**: Can the C-state agreement, membership service, and temporal encapsulation concepts be implemented as a TSN profile at 1 Gbit/s, providing integrated fault tolerance for modern zonal automotive architectures?
- **Formal verification of mode change and startup**: The safety of TTP's cold-start, re-integration, and mode-change protocols under all asynchronous fault patterns has not been fully formally verified. Model checking (e.g., with TLA+ or SPIN) of the full state space under concurrent faults is an open task.
- **Energy-aware TDMA**: Can TTP's deterministic slot structure be exploited for fine-grained per-slot power gating in battery-constrained embedded nodes, and what is the achievable energy reduction relative to CAN?

References

[1] H. Kopetz and G. Grünsteidl, "TTP – A Time-Triggered Protocol for Fault-Tolerant Real-Time Systems," in *Proc. 23rd Int. Symp. Fault-Tolerant Computing (FTCS-23)*, pp. 524–533, Toulouse, France, June 1993.

[2] H. Kopetz, "Event-Triggered versus Time-Triggered Real-Time Systems," in *Proc. Int. Workshop on Operating Systems of the 90s and Beyond*, LNCS vol. 563, pp. 87–101, Springer-Verlag, 1991.

[3] SAE Handbook, Vol. 2. Society of Automotive Engineers, PA, USA, 1992.

[4] FlexRay Consortium, *FlexRay Communications System – Protocol Specification*, Version 2.1, 2005.

[5] L. Lamport, R. Shostak, and M. Pease, "The Byzantine Generals Problem," *ACM Trans. on Programming Languages and Systems*, vol. 4, no. 3, pp. 382–401, 1982.

[6] S. Vestal, "Preemptive Scheduling of Multi-Criticality Systems with Varying Degrees of Execution Time Assurance," in *Proc. 28th IEEE Real-Time Systems Symp. (RTSS)*, pp. 239–243, Tucson, AZ, USA, 2007.

[7] IEEE Standards Association, "IEEE 802.1Qbv – Enhancements for Scheduled Traffic," IEEE Standard, 2015.